

## CSCI 558L – Laboratory Assignment #7: Internet Router Experience

Instructor: Young H. Cho

T.A.: Siddharth Bhargav

**Due date: September 25, 2014 at 11:55pm**

This project will let you gain some experience setting up static routing on a small network. The NS file you should use is attached as listing 1. We will use a network of 4 hosts and 3 routers. Note the 2<sup>nd</sup> to last line of the NS file:

```
$ns rtrproto Manual
```

Normally this is set to static, and the DETER system sets the kernel routing tables for your all of your hosts so they can talk to each other. Here we set it to manual, which means the hosts and routers do not know about the network topology. Your task is to properly configure the routers so the 4 hosts can exchange packets.

We will also be using apart of an open source IP routing software called Quagga (<http://www.quagga.net/>). Quagga implements several routing protocols (RIP, OSPF, BGP, and etc.) and provides a 'Cisco-like command-line interface for configuring and interacting with the router software. Quagga runs on top of most modern UNIX-like operating systems and is made up of several encapsulated software pieces called daemons. Quagga does not actually route packets, it uses information from the implemented protocols to update the kernel routing table. The piece of Quagga we'll use for this assignment is called zebra. Zebra is the management daemon responsible for updating the kernel routing table. When used with the other routing daemons (ospfd, ripd, bgpd) zebra takes updates from these daemons and applies them to the kernel routing table. However, zebra can also be used for managing static routes.

This homework will require you to build and use the Quagga software. This handout will not provide extremely detailed instructions, just enough to get you started. This way we hope you'll learn the most. This assignment uses 10 total DETER nodes, so you MUST swap out or terminate your experiments when you are not using them!

- Begin by starting the experiment with the NS file shown in listing 1. Can node1 ping node3? Can node1 ping rtr1?
- Where does the default route on node1 point? Where does the default route on node3 point? All of the nodes in DETER are connected to the control network (with IP address 192.168.X.X), we need to pretend this connection doesn't exist. Change the default route for the 4 hosts to their respective router (rtr1 or rtr3 respectively). Use the route command to do so. This should be the only change necessary on the 4 hosts. Include a print out of the kernel routing table before and after you make this change. Important: before you change the default route, enter the following command:

```
sudo route add -host 192.168.253.1 gw 192.168.1.254
```

Otherwise once the default route is changed, you won't be able to connect to your node from users.isi.deterlab.net .

- Download the Quagga source code to your users.deterlab.net account. On each router node use the mkextrafs.pl script to make a filesystem at /mnt. Copy the Quagga source to /mnt and unpack it (also on each node).
- Build the Quagga software by executing “./configure --enable-user=root --enable-group=root” and make in the Quagga source directory. Do not install the software using make install. Do this on all three router nodes.
- You should now have the Quagga software built. Create a configuration file for zebra in /mnt. Call it zebra.conf. Below is an example:

```
!  
! Example zebra configuration  
!  
hostname zebrad  
password zebra  
!  
interface eth5  
ip address 10.10.1.1/24  
!  
interface eth7  
ip address 10.1.0.3/24  
!  
interface eth9  
ip address 10.10.2.2/24  
!  
interface lo  
ip forwarding  
!  
log stdout  
line vty  
!
```

In this file we need to tell zebra about the interfaces we're going to use. Note we ignore the control network. You will probably have to change the names of the interfaces to match the specific machine you are assigned by DETER.

- Perform a similar procedure on rtr2 and rtr3. You will need to create the configuration file yourself. Once you have it correct, include the configuration files in your report.

- Now we need to start zebra on the routers. Change into the zebra subdirectory in the Quagga source code. Execute the following command:

```
sudo ./zebra -d -u root -f /mnt/conf/zebra.conf
```

- Now the fun begins! Lets ignore rtr2 for now. On rtr1 and rtr3 we need to set static routes that tell the routers where to find the subnet behind the other router. So we need to tell rtr1 about the 10.1.2.0/24 subnet and we need to tell rtr3 about the 10.1.0.0/24 subnet. Look at the Quagga documentation (it's online) and figure out how to do this. Start by connecting to the zebra daemon:

```
telnet localhost 2601
```

To enter commands that change the configuration use enable followed by configure terminal. It should take only one command on each router ( rtr1 and rtr3 ) to enable the two LANs to exchange packets.

- Use ping to verify you can now exchange packets between the LANs. About how much RTT latency is there? Does this make sense?
- Include in your report the commands you used to configure the routes in zebra, a traceroute between node1 and node4, the routing tables in the routers, and an explanation of what you did and why it works.
- Now let's look at rtr2. Can rtr2 ping any of the nodes in the LANs? Why not? Add two routes to rtr2 such that rtr2 can ping nodes in the LANs. Include the commands you used and the routing table in rtr2 in your report.
- Now go back and add a secondary route through rtr2 on rtr1 and rtr3. When you do show ip route on either rtr1 or rtr3 you should see two lines in the routing table for one entry. Similar to:

```
S>* 10.1.0.0/24 [1/0] via 10.10.1.1, eth7
                        via 10.10.3.2, eth10
```

- Does adding these secondary routes change the path a packet takes from LAN to LAN? (i.e. did the traceroute from node1 to node4 change?)
- Now set the loss ratio on link0 (the link between rtr1 and rtr3) to 1.0. Can you still ping between the LANs? What does traceroute show?
- Now change the metric (distance) for the routes in rtr1 and rtr3. Set the route through rtr2 to have a metric of 1 and change the other route (through rtr1 or rtr3 respectively) to have a metric

of 5. When you have it correct, part of the output from show ip route on rtr1 should look similar to:

```
S>* 10.1.2.0/24 [1/0] via 10.10.2.1, eth9
S 10.1.2.0/24 [5/0] via 10.10.1.2, eth5
```

- Can you now ping from LAN to LAN? About how much delay is there now? Does this make sense? How about traceroute? Your traceroute should have one hop that does not return (i.e. you get \* \* \* ). Explain why changing the metric 'fixed' the network? Explain why traceroute still has an error?

## Discussion

- What does setting the link0 loss to 1.0 simulate?
- How long do you think it would take you to change or update the routes manually if you worked at an ISP and a link between routers went down? Is this a problem? What if you had to change the routes on 5 routers, 10 routers? Does manually updating the routing tables scale well?
- How would protocols like OSPF solve this problem?

### Listing 1: Lab #5 NS file

```
set ns [new Simulator]
source tb_compat.tcl
# Nodes
set rtr1 [$ns node]
tb-set-node-os $rtr1 Ubuntu1004-STD
set rtr2 [$ns node]
tb-set-node-os $rtr2 Ubuntu1004-STD
set rtr3 [$ns node]
tb-set-node-os $rtr3 Ubuntu1004-STD
set node1 [$ns node]
tb-set-node-os $node1 Ubuntu1004-STD
set node2 [$ns node]
tb-set-node-os $node2 Ubuntu1004-STD
set node3 [$ns node]
tb-set-node-os $node3 Ubuntu1004-STD
set node4 [$ns node]
tb-set-node-os $node4 Ubuntu1004-STD

# Links
set link0 [$ns duplex-link $rtr1 $rtr3 1000000.0kb 5.0ms DropTail]
tb-set-ip-link $rtr1 $link0 10.10.1.1
tb-set-ip-link $rtr3 $link0 10.10.1.2
set link1 [$ns duplex-link $rtr2 $rtr3 1000000.0kb 5.0ms DropTail]
tb-set-ip-link $rtr3 $link1 10.10.3.1
tb-set-ip-link $rtr2 $link1 10.10.3.2
set link2 [$ns duplex-link $rtr1 $rtr2 1000000.0kb 5.0ms DropTail]
tb-set-ip-link $rtr2 $link2 10.10.2.1
tb-set-ip-link $rtr1 $link2 10.10.2.2

# Lans
set lan0 [$ns make-lan "$node1 $rtr1 $node2" 1000000.0kb 0.0ms]
tb-set-ip-lan $node1 $lan0 10.1.0.1
tb-set-node-lan-bandwidth $node1 $lan0 1000000.0kb
tb-set-ip-lan $rtr1 $lan0 10.1.0.3
tb-set-node-lan-bandwidth $rtr1 $lan0 1000000.0kb
tb-set-ip-lan $node2 $lan0 10.1.0.2
tb-set-node-lan-bandwidth $node2 $lan0 1000000.0kb
set lan1 [$ns make-lan "$rtr3 $node3 $node4" 1000000.0kb 0.0ms]
tb-set-ip-lan $rtr3 $lan1 10.1.2.1
tb-set-node-lan-bandwidth $rtr3 $lan1 1000000.0kb
tb-set-ip-lan $node3 $lan1 10.1.2.3
tb-set-node-lan-bandwidth $node3 $lan1 1000000.0kb
tb-set-ip-lan $node4 $lan1 10.1.2.4
tb-set-node-lan-bandwidth $node4 $lan1 1000000.0kb

$ns rtproto Manual
$ns run
```